Badge-2 Lab-3 [Linear Regression]

Out date: 20 July 2022
Due date: 24 July 2022 @ 11:59pm

Objective: To apply multi-linear regression to predict a quantitative target variable.

Data:
Horizontal drilling which provides additional exposure of the reservoir to the wellbore (lateral length) is one of the primary drivers that made economic production of oil and gas from the tight shale reservoirs in the US a success, in conjunction with stimulation using hydraulic fracturing.

Objective of the exercise is to understand effect of lateral length of a horizontal well on shale gas production in some of the US shale plays using multiple linear regression analysis. While lateral length would be one of the primary predictors, other variables considered are vertical section length of the horizontal well, proppant volume and fluid volume pumped for stimulating the horizontal section of the well using hydraulic fracturing and type of shale plays.

Data source: National Oil and Gas Gateway (http://www.noggateway.org/explore). From participating states that contribute to this data source, Arkansas and Colorado were identified as states of interest for the project owing to the level of shale gas drilling and production activity in these states over the period 2008-2018 and data availability on Fluids and Proppants.

Variables of interest:

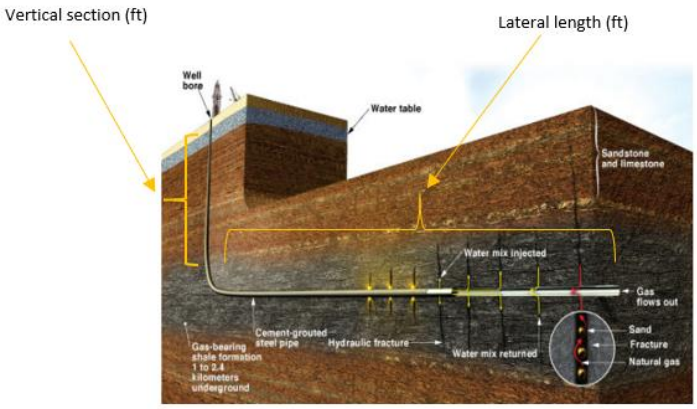

Figure 1: Horizontal well schematic (from The Academy of Medicine, 2017)

Referring to the Figure-1 above, the following were identified as variables of interest after reviewing the dataset:

Response (Target) variable: Max_Gas, Maximum Gas Production (Million Standard Cubic Feet, MSCF)
Since shale gas wells have their peak production in the first 2-3 years after the well is drilled, stimulated and put on production, maximum annual gas production is selected as the best response variable to predict using the predictor variables identified below

Predictor variables:
Lat_Len, Lateral length of the horizontal section (feet)

Measured Depth, Total Depth of the well (**feet**)

Bottom and **Top** Perforation depth (feet)

Wellhead **Latidude** and **Longitude** (decimal degrees), surface location of the well

Fluid, Total amount of fracturing fluid (typically water) pumped to create the hydraulic fractures (Barrels)

Proppant, Total amount of proppant (typically sand) pumped along with the fracturing fluid as a slurry to keep the hydraulically created fractures open (Pounds)

ShalePlay (a dummy variable, 0- Fayetteville, Arkansas Shale gas play and 1- Mancos, Colorado Shale gas play)--**STATE**

Drilling Orientation (a dummy variable)

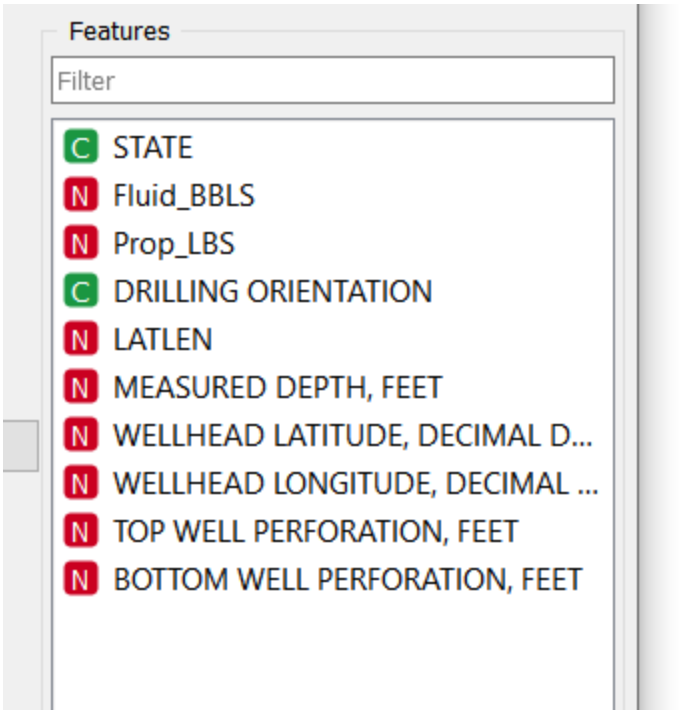

Data: For this lab, please download *production.csv* and *gas_prod_pred_start.ows* to your local folder.

Lab Instructions

1. Open the *gas_prod_pred_start.ows* file in Orange. Your pipeline should look as below. Confirm the *data.xlsx* dataset is loaded by opening the **File** widget.

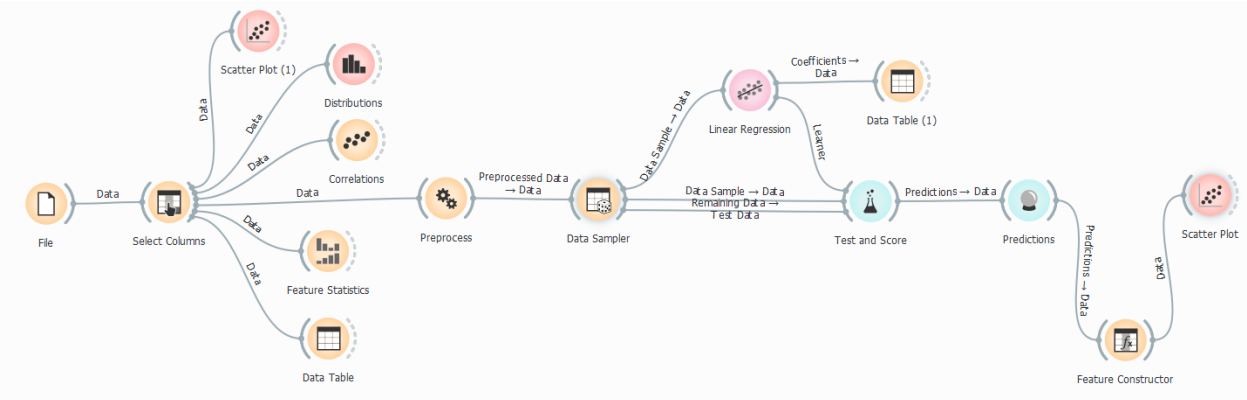

2. Inspect the **File** widget and complete the following table.

Total instances	6310
Dimensionality of the data set	Rows: 6310 Columns: 14
Predictors (features)	10

Target variable	MAXGAS

3. Use the **Data Table, Feature Statistics, Distribution, Scatter Plot & Correlations and Scatter Plot** widgets to complete the table below:

Min and Max of predictor variables	
Do you notice any data inconsistency?	Yes, like depth -592 (depth cannot be negative, below depth)
Distribution of the predictor variables:	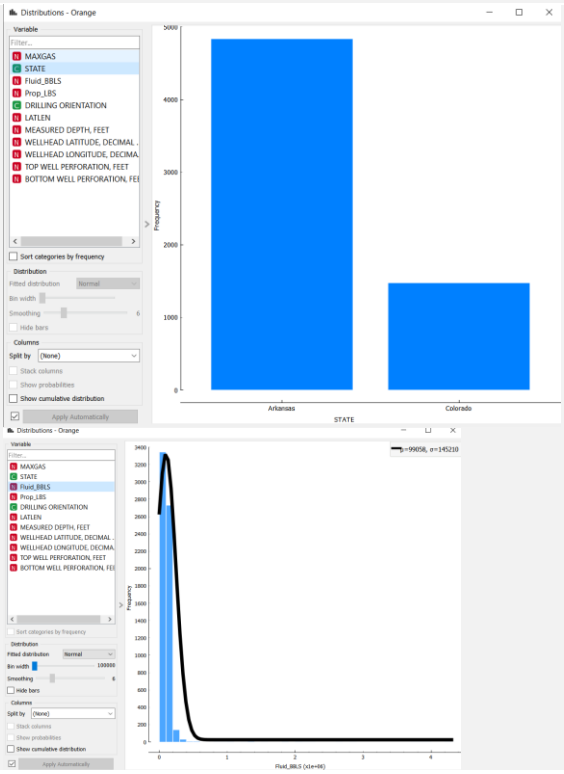

Feature Statistics - Orange

Name	Distribution	Mean	Median	Dispersion	Min.	Max.	Missing
Prop_LBS		2888474.46	3233356.50	0.65	0	9964445	0 (0%)
MAXGAS		377094.02	352435.50	0.63	264	2594846	0 (0%)
Fluid_BBLS		99058	95599	1.46591	0.00	4.25904e+06	36 (0%)
MEASURED DEPTH, FEET		8495.44	8491	0.25	-592	17645	0 (0%)
BOTTOM WELL PERFORATI... FEET		8297.54	8322	0.27	1095	17457	0 (0%)
TOP WELL PERFORATI... FEET		4775.68	4541	0.35	244	13514	0 (0%)
LATLEN		3521.86	3922	0.53	2	9605	0 (0%)
WELLHEAD LATITUDE, DECIMAL DEGREES		36.377856	35.4593	0.049306	35.0236	40.9722	1 (0%)
WELLHEAD LONGITUDE, DECIMAL DEGREES		-95.862870	-92.5018	-0.065459	-108.9530	-91.5073	1 (0%)
STATE			Arkansas	0.544			0 (0%)
DRILLING ORIENTATION			Horizontal	0.802			0 (0%)

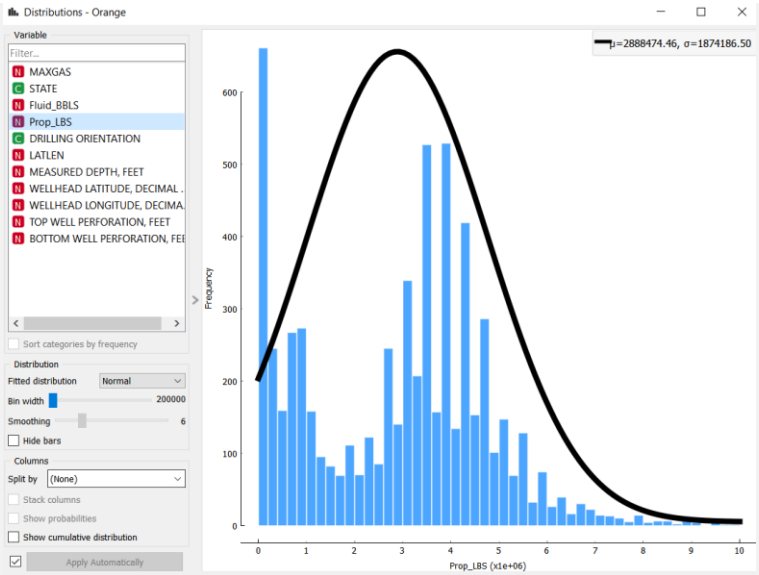
Distributions - Orange
Variable
Filter...
MAXGAS
STATE
Fluid_BBLS
Prop_LBS
DRILLING ORIENTATION
LATLEN
MEASURED DEPTH, FEET
WELLHEAD LATITUDE, DECIMAL .
WELLHEAD LONGITUDE, DECIMA.
TOP WELL PERFORATION, FEET
BOTTOM WELL PERFORATION, FEE
Sort categories by frequency
Distribution
Fitted distribution
Normal
Bin width
200000
Smoothing
6
Hide bars
Columns
Split by
(None)
Stack columns
Show probabilities
Show cumulative distribution
Apply Automatically
μ=2888474.46, σ=1874186.50
Frequency
Prop_LBS (x1e+06)

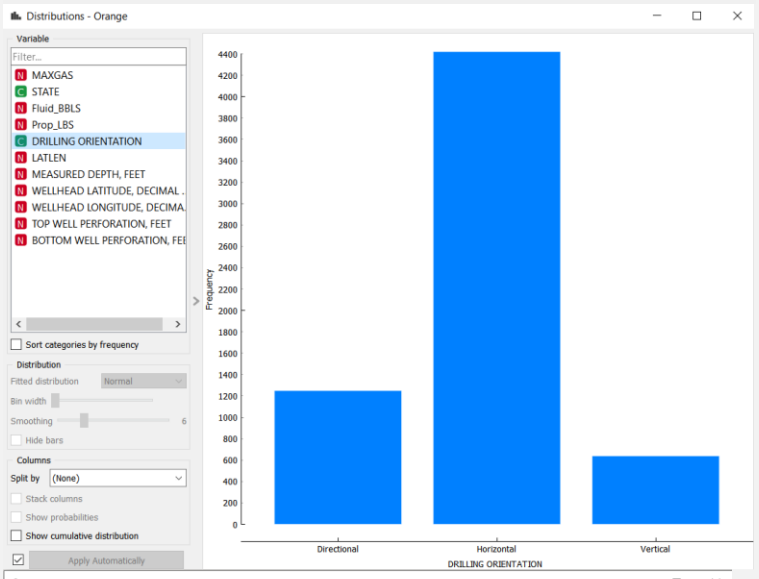
Distributions - Orange
Variable
Filter...
MAXGAS
STATE
Fluid_BBLS
Prop_LBS
DRILLING ORIENTATION
LATLEN
MEASURED DEPTH, FEET
WELLHEAD LATITUDE, DECIMAL .
WELLHEAD LONGITUDE, DECIMA.
TOP WELL PERFORATION, FEET
BOTTOM WELL PERFORATION, FEE
Sort categories by frequency
Distribution
Fitted distribution
Normal
Bin width
Smoothing
6
Hide bars
Columns
Split by
(None)
Stack columns
Show probabilities
Show cumulative distribution
Apply Automatically
Frequency
Directional
Horizontal
Vertical
DRILLING ORIENTATION

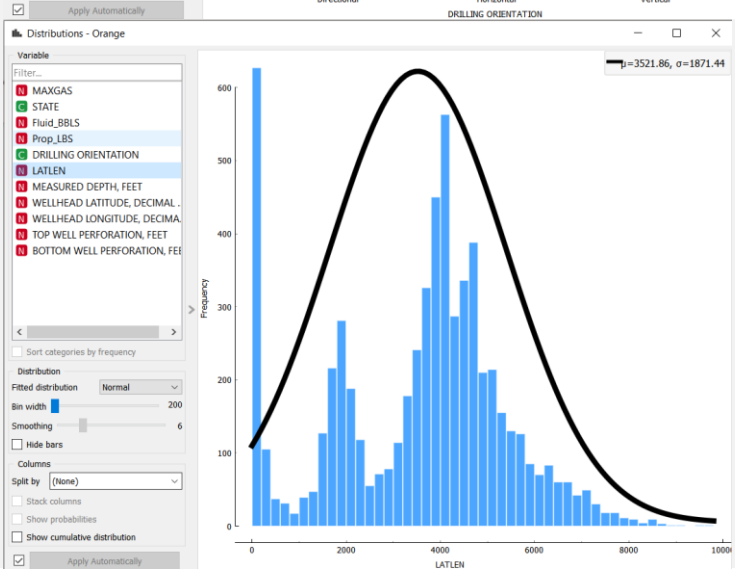
Distributions - Orange
Variable
Filter...
MAXGAS
STATE
Fluid_BBLS
Prop_LBS
DRILLING ORIENTATION
LATLEN
MEASURED DEPTH, FEET
WELLHEAD LATITUDE, DECIMAL .
WELLHEAD LONGITUDE, DECIMA.
TOP WELL PERFORATION, FEET
BOTTOM WELL PERFORATION, FEE
Sort categories by frequency
Distribution
Fitted distribution
Normal
Bin width
200
Smoothing
6
Hide bars
Columns
Split by
(None)
Stack columns
Show probabilities
Show cumulative distribution
Apply Automatically
μ=3521.86, σ=1871.44
Frequency
LATLEN

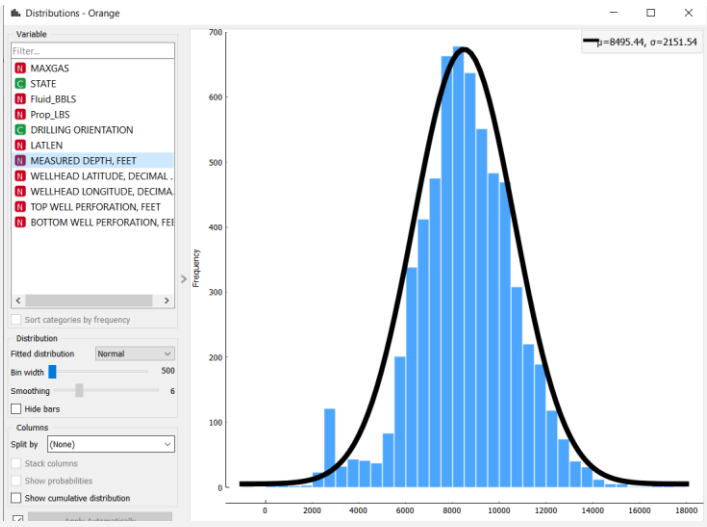
Distributions - Orange
Variable
Filter...
MAXGAS
STATE
Fluid_BBLS
Prop_LBS
DRILLING ORIENTATION
LATLEN
MEASURED DEPTH, FEET
WELLHEAD LATITUDE, DECIMAL .
WELLHEAD LONGITUDE, DECIMA.
TOP WELL PERFORATION, FEET
BOTTOM WELL PERFORATION, FEE
Sort categories by frequency
Distribution
Fitted distribution
Normal
Bin width
500
Smoothing
6
Hide bars
Columns
Split by
(None)
Stack columns
Show probabilities
Show cumulative distribution
μ=8495.44, σ=2151.54
Frequency
0
100
200
300
400
500
600
700
0
2000
4000
6000
8000
10000
12000
14000
16000
18000

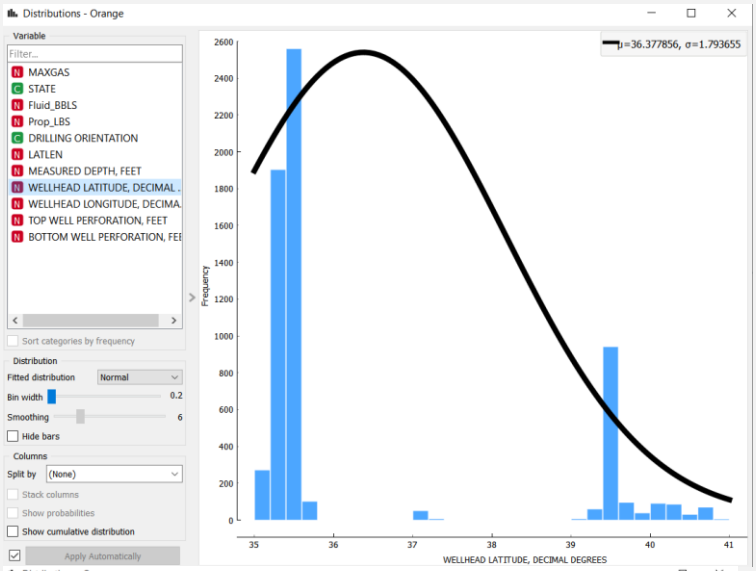
Distributions - Orange
Variable
Filter...
MAXGAS
STATE
Fluid_BBLS
Prop_LBS
DRILLING ORIENTATION
LATLEN
MEASURED DEPTH, FEET
WELLHEAD LATITUDE, DECIMAL .
WELLHEAD LONGITUDE, DECIMA.
TOP WELL PERFORATION, FEET
BOTTOM WELL PERFORATION, FEE
Sort categories by frequency
Distribution
Fitted distribution
Normal
Bin width
0.2
Smoothing
6
Hide bars
Columns
Split by
(None)
Stack columns
Show probabilities
Show cumulative distribution
Apply Automatically
μ=36.377856, σ=1.793655
Frequency
0
200
400
600
800
1000
1200
1400
1600
1800
2000
2200
2400
2600
35
36
37
38
39
40
41
WELLHEAD LATITUDE, DECIMAL DEGREES

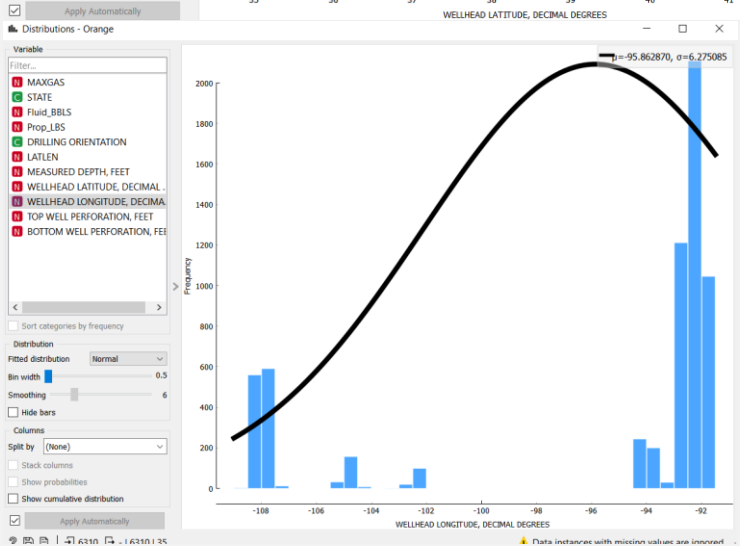
Distributions - Orange
Variable
Filter...
MAXGAS
STATE
Fluid_BBLS
Prop_LBS
DRILLING ORIENTATION
LATLEN
MEASURED DEPTH, FEET
WELLHEAD LATITUDE, DECIMAL .
WELLHEAD LONGITUDE, DECIMA.
TOP WELL PERFORATION, FEET
BOTTOM WELL PERFORATION, FEE
Sort categories by frequency
Distribution
Fitted distribution
Normal
Bin width
0.5
Smoothing
6
Hide bars
Columns
Split by
(None)
Stack columns
Show probabilities
Show cumulative distribution
Apply Automatically
μ=-95.862870, σ=6.275085
Frequency
0
200
400
600
800
1000
1200
1400
1600
1800
2000
-108
-106
-104
-102
-100
-98
-96
-94
-92
WELLHEAD LONGITUDE, DECIMAL DEGREES

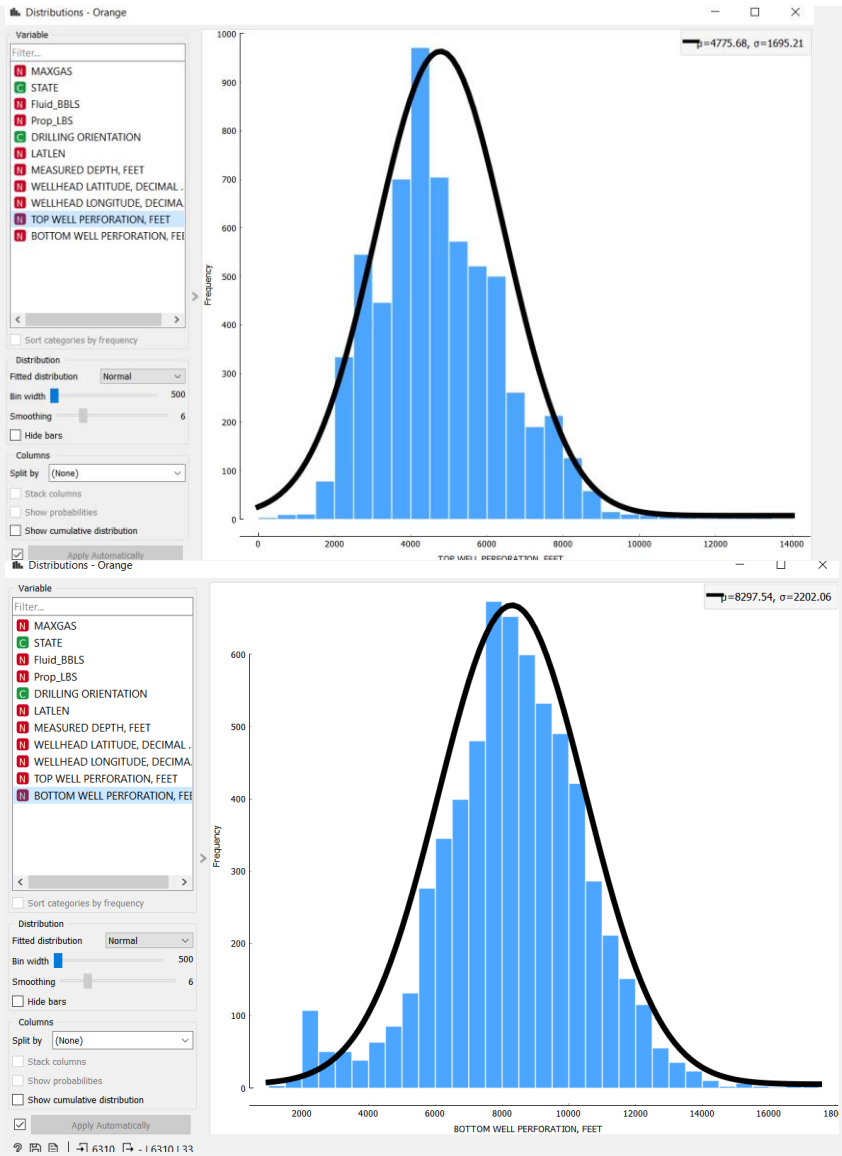

Min, Max and distribution of the target variable, Max_Gas

Min:264, Max: 2594846, skewed distribution

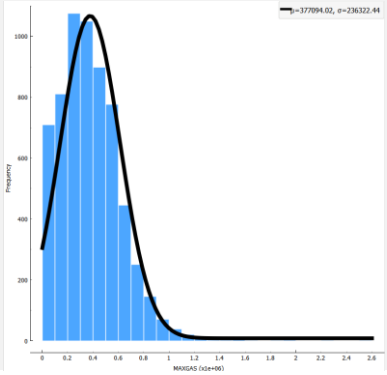

% of wells from the two states (shale plays) and drilling orientation	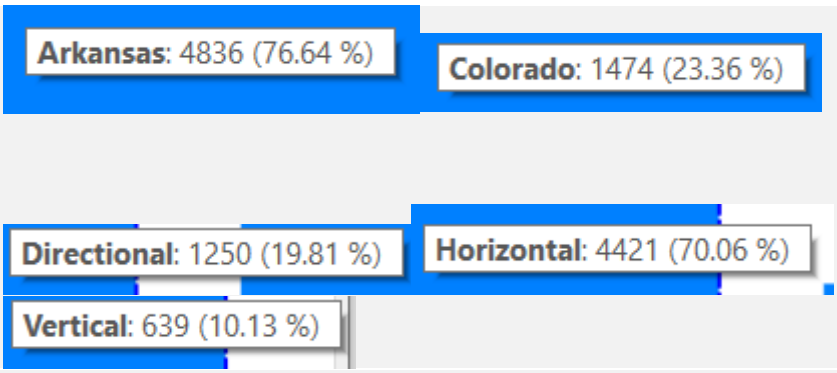
Any missing values?	yes
From the **Scatter Plot**, what are the predictors that you believe would have the highest impact on predicting Max_Gas?	LATLEN (60% strongest correlation), Prop_LBS, BOTTOM WELL… 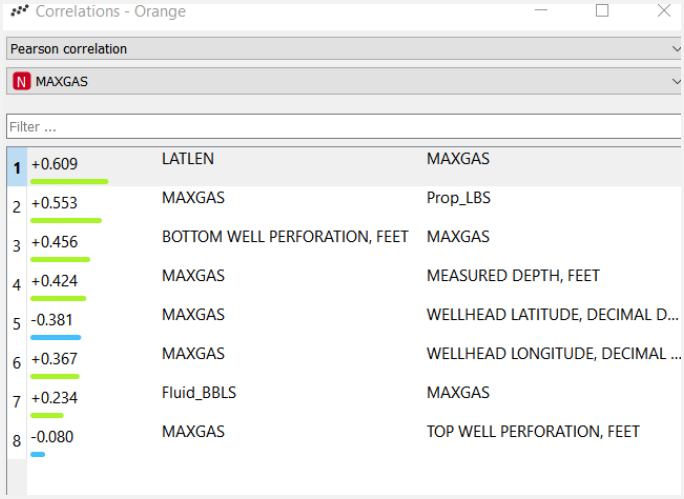
How would you describe the relationships for these predictors with the target?	Increase of predictors usually mean more gas productions (MAXGAS)
Do you notice outliers in the data?	Yes, for example: X: LATLEN, Y: MAXGAS, and has an outlier close to 2.6

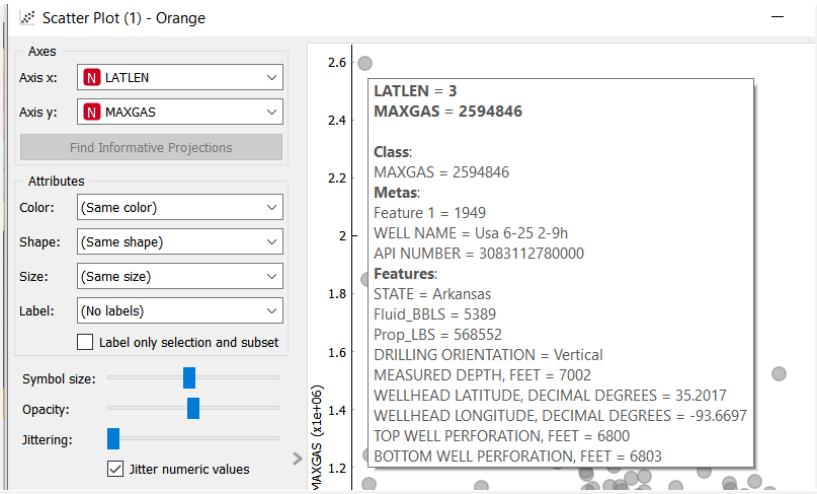

From the **correlations** widget, what do you think about the correlation between predictors?

LATLEN 60% strongest correlation with MAXGAS, second is Prop_LBS (55%)

some have really strong correlations with each other (LATLEn and Prop_LBS 88%)

Each predictors have their strong correlations, expected to be independent

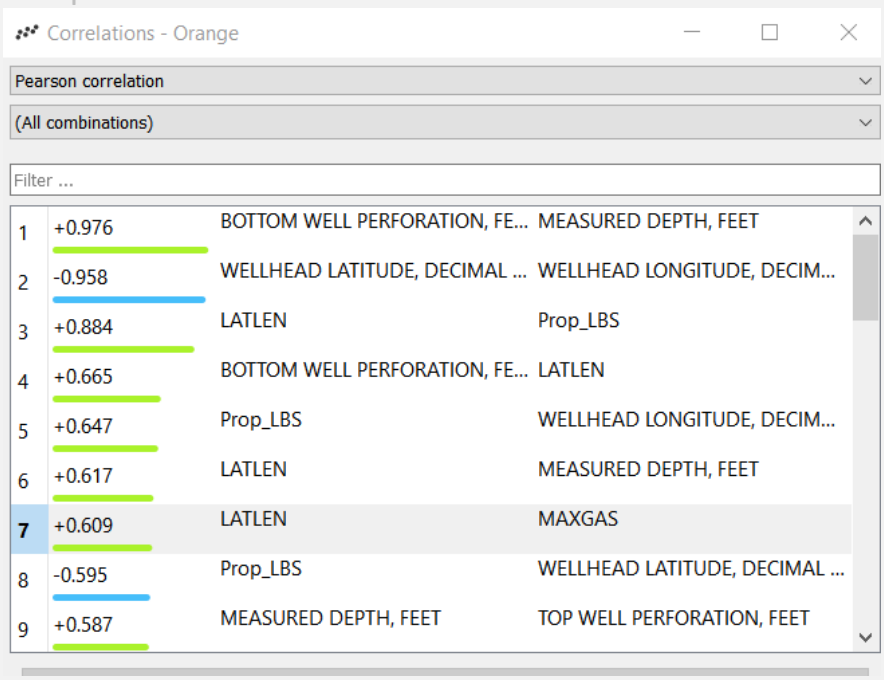

1	+0.976	BOTTOM WELL PERFORATION, FE...	MEASURED DEPTH, FEET
2	-0.958	WELLHEAD LATITUDE, DECIMAL ...	WELLHEAD LONGITUDE, DECIM...
3	+0.884	LATLEN	Prop_LBS
4	+0.665	BOTTOM WELL PERFORATION, FE...	LATLEN
5	+0.647	Prop_LBS	WELLHEAD LONGITUDE, DECIM...
6	+0.617	LATLEN	MEASURED DEPTH, FEET
7	+0.609	LATLEN	MAXGAS
8	-0.595	Prop_LBS	WELLHEAD LATITUDE, DECIMAL ...
9	+0.587	MEASURED DEPTH, FEET	TOP WELL PERFORATION, FEET

4. Address data inconsistency by adding **Select Rows** widget and filtering them out.

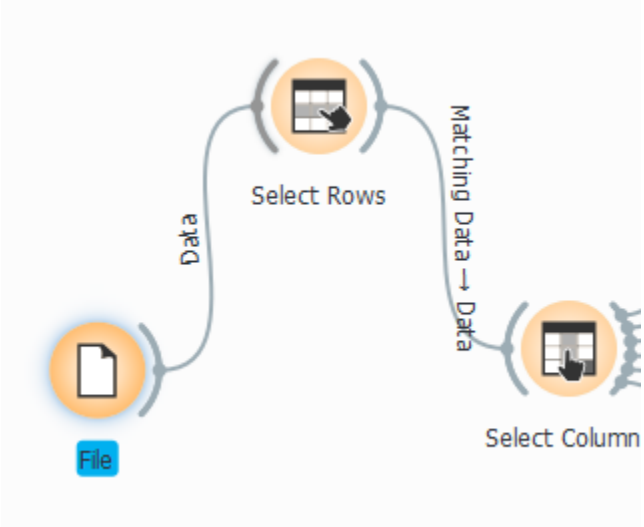

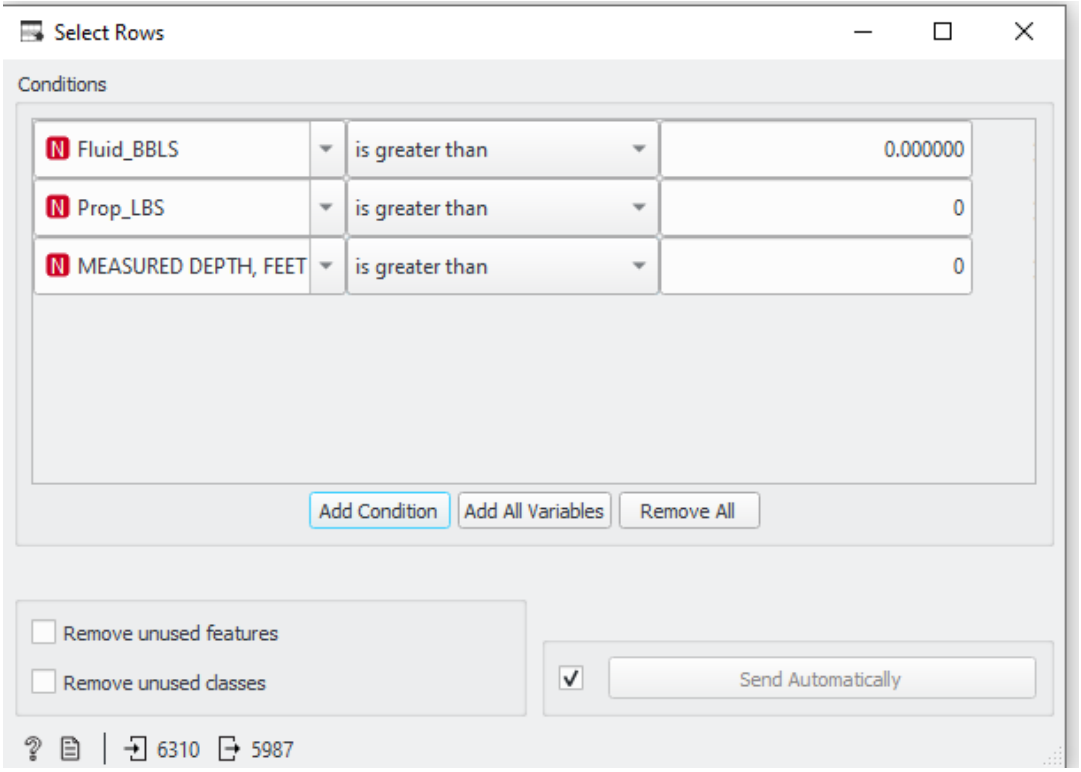

How many instances were filtered out in the above step? Eliminated 323 rows

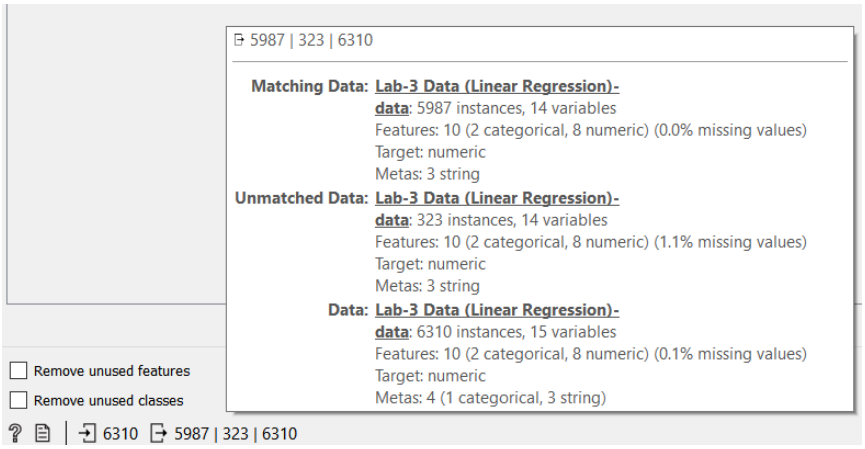

5. Open **Select Columns** widget and confirm selection as below:

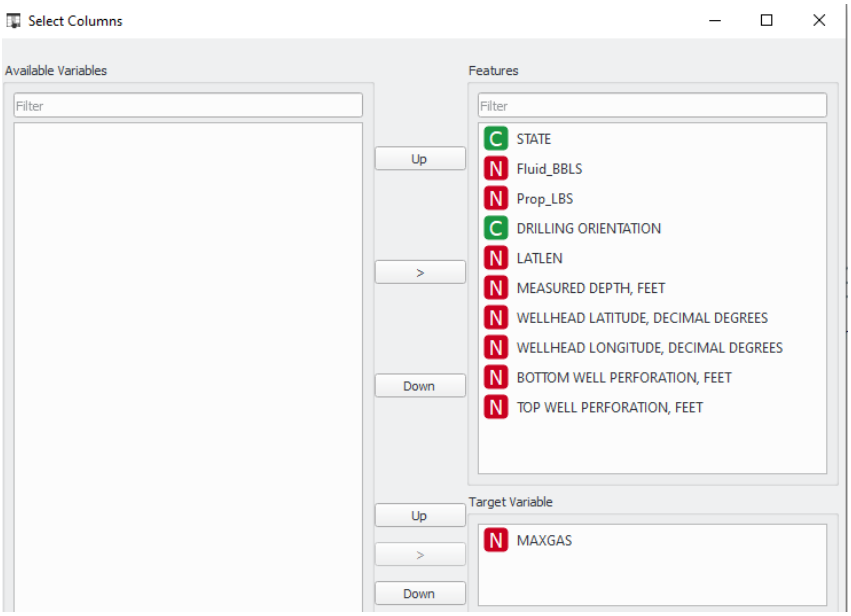

6. Inspect the **Preprocess** and **Data Sampler** widgets.

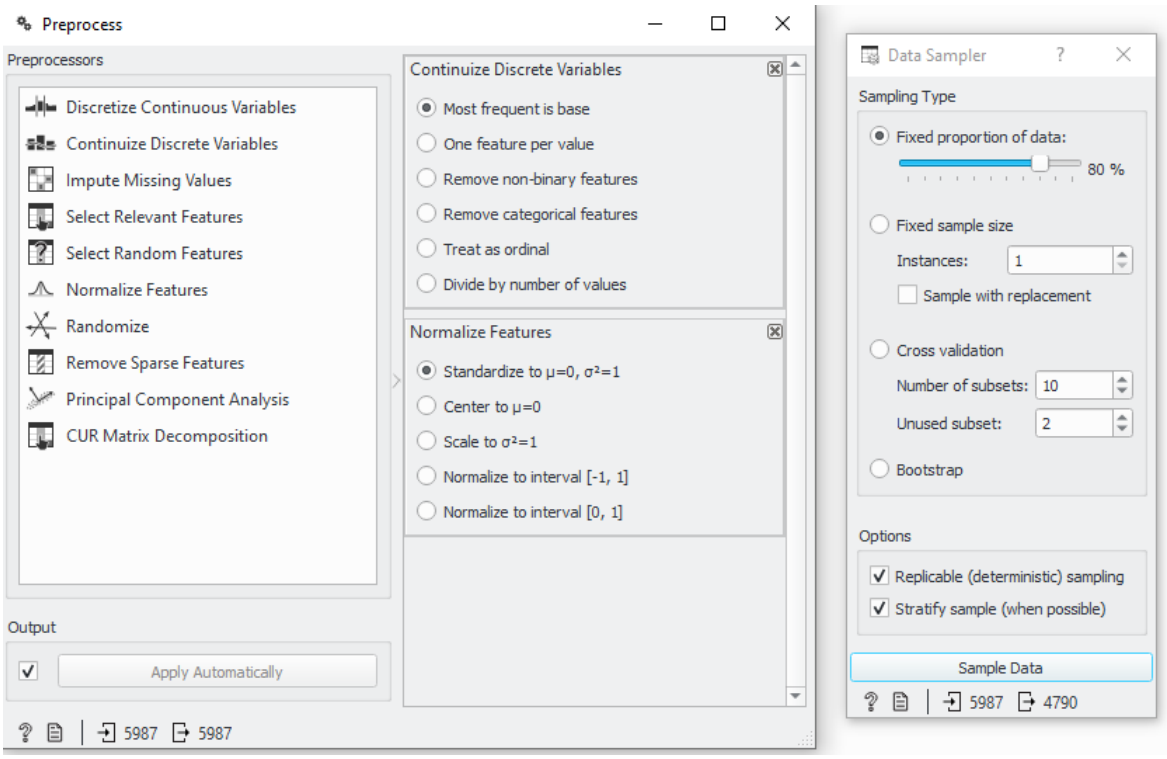

7. Open the **Linear Regression** widget and inspect the same:

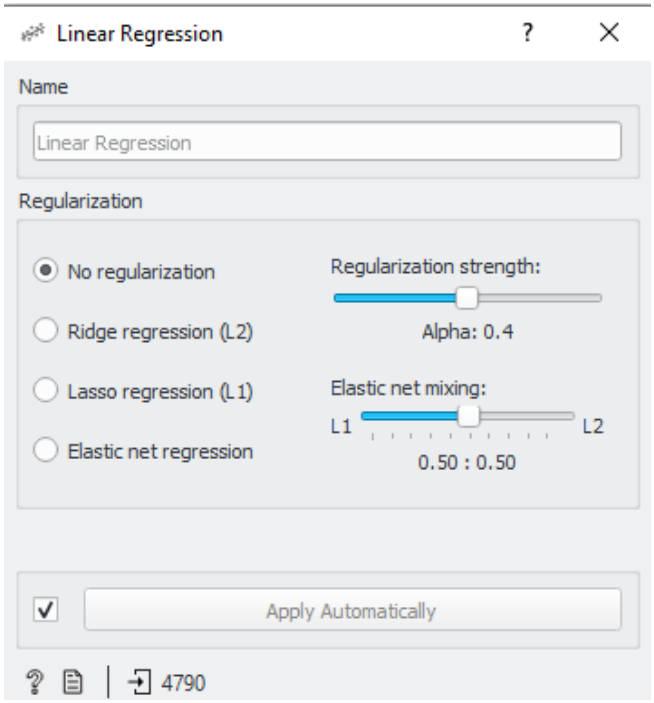

8. Open **Data Table** widget connected to the **Linear Regression** widget and inspect the parameter coefficients.

	name	coef
1	intercept	380225
7	LATLEN	65041.7
11	BOTTOM WELL PERFORATION, FEET	57315.5
4	Prop_LBS	44794.3
3	Fluid_BBLS	6951.74
12	TOP WELL PERFORATION, FEET	3024.7
6	DRILLING ORIENTATION=Vertical	-35943.4
8	MEASURED DEPTH, FEET	-43826
9	WELLHEAD LATITUDE, DECIMAL DEGREES	-47177.5
5	DRILLING ORIENTATION=Directional	-116990
2	STATE=Colorado	-181543
10	WELLHEAD LONGITUDE, DECIMAL DEGREES	-311859

What can you say about the relationships between the different predictors and the response variable (Max_Gas)?	They influence or relate some way (LATLEN most important variable and strongest influence); vertical wells will have a degraded performance while horizontal perform better; measured depth is actually having negative influence to the production

9. Open **Test and Score** widget and complete the table below:

What are the evaluation metrics you see?	Evaluation Results Model | Train time [s] | Test time [s] | RMSE | MAE | R2 Linear Regression | 0.122 | 0.021 | 181848.288 | 132069.860 | 0.407

Model	RMSE	MAE	R2
Linear Regression (cross validation , 10 folds	181848.288	132069.860	0.407
Linear Regression (Test on test data)	182165.880	131290.348	0.396

How would you interpret this model based on R2 and cross validation?	Decreased percentage, and given R2 is 40 percent vs 100 (or high number), it is not a perfect model

Open the **Scatter Plot** widget connected to **Predictions widget.** Select Max_Gas as X-axis and Linear Regression as Y-axis.

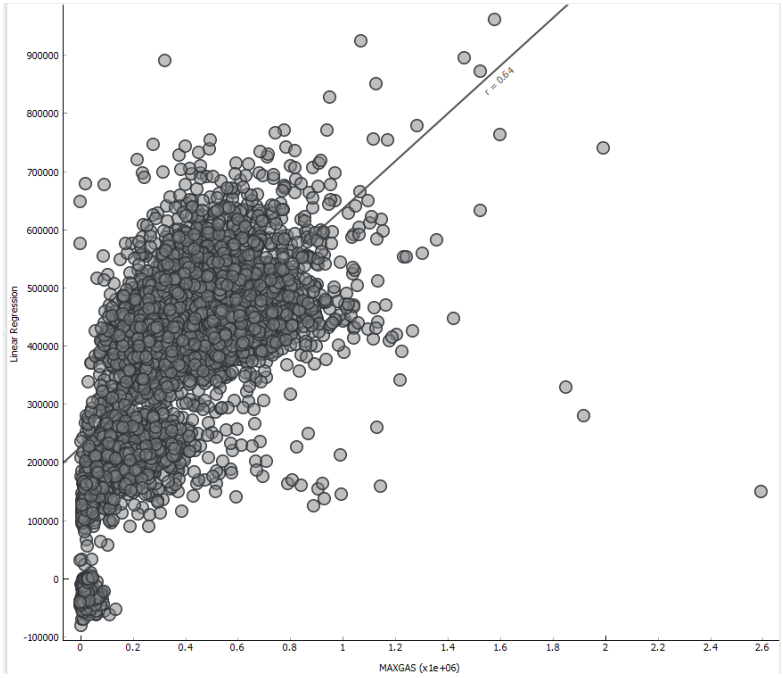

What do you think about the prediction quality?	Not too perfect since not all the plots are lined to the r and likely not deploy the model, but still wouldn't discard this to see what can be improved

10. As part of model diagnostics in linear regression, it is standard practice to visualize the residual (error) plot and look for patterns to understand model validity.
    Residual variable is constructed as shown below using the **Feature Constructor** widget:

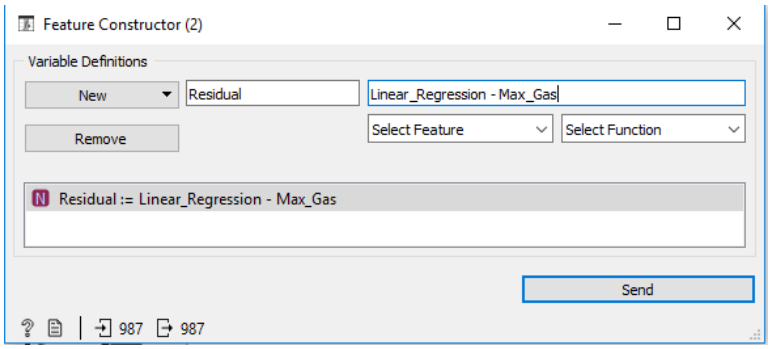

11. Open **Scatter Plot** widget and select Linear Regression (X axis)-(Predicted Max_Gas from the regression model) and Residual (Y axis). Plot would look as below:

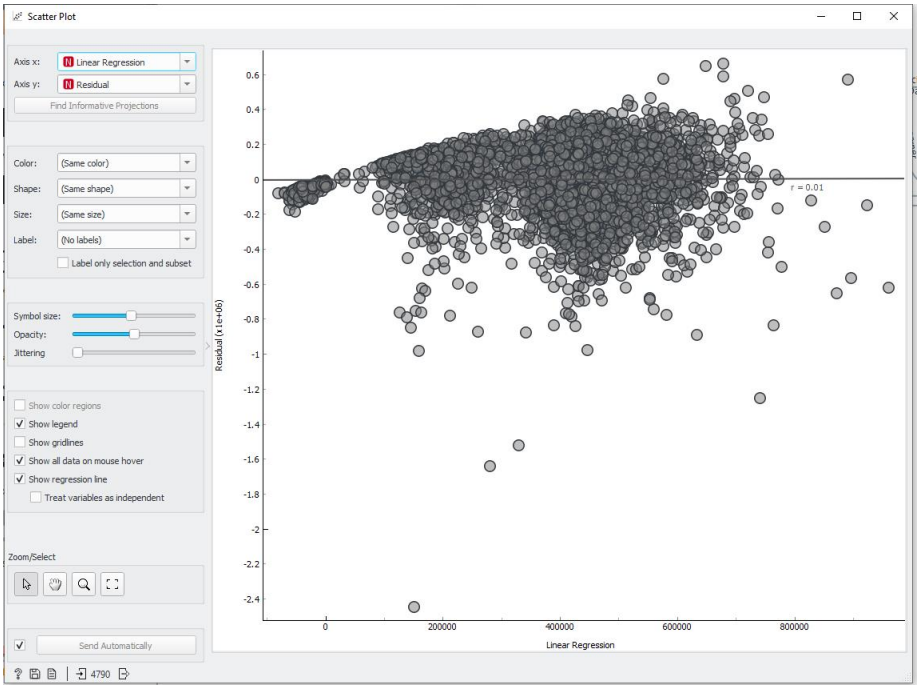

Do you see a pattern?	Center line running through 0, thus a corn/megaphone pattern
What does this indicate?	Average out to mean; when you train a linear regression model, for it to be valid, the error should have a normal distribution with a mean of 0
What can you do to address this issue?	Indicates that tho it looks okay, it is not quite valid Address by apply transformation (look at log or square root of variable) to your target variable (changes distribution of target to make it valid more likely)

12. Let us transform the target variable Max_Gas using Log. Add another **Feature Constructor** widget near the **File widget.** Construct Log transformation as shown below:

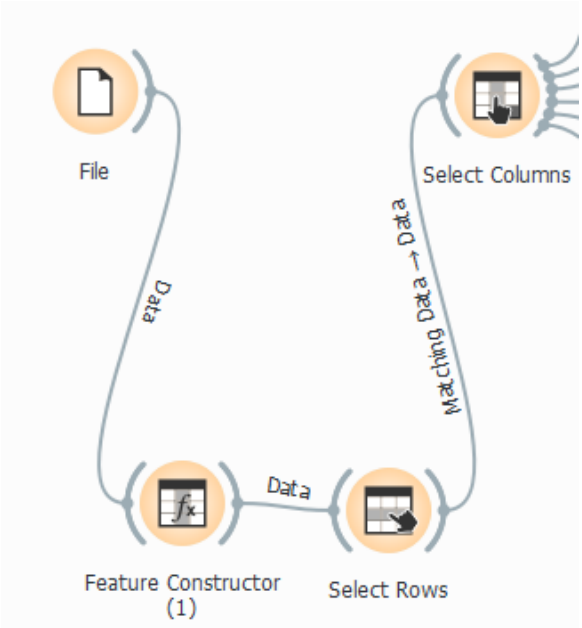

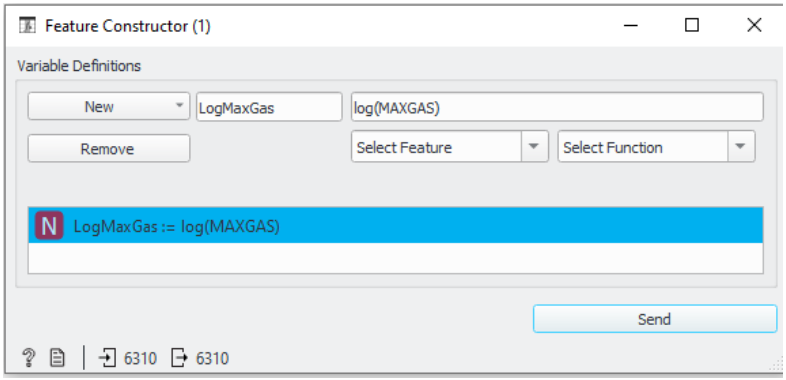

Modify **Select Columns** widget as below:

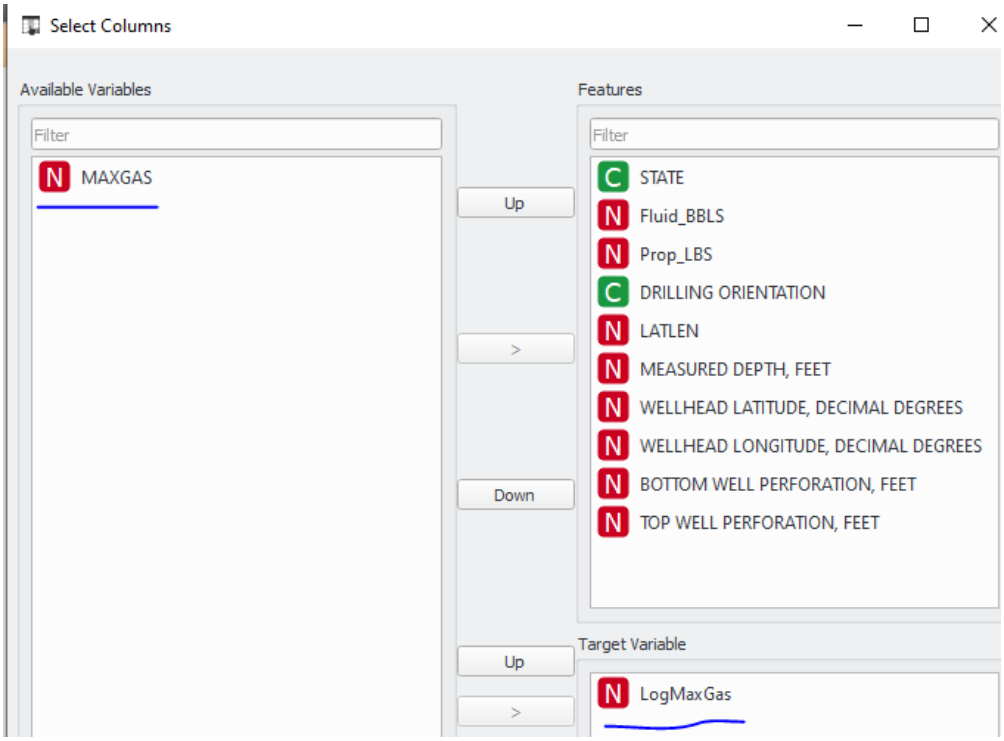

Update the second Feature constructor (farthest right) to calculate the residual correctly:

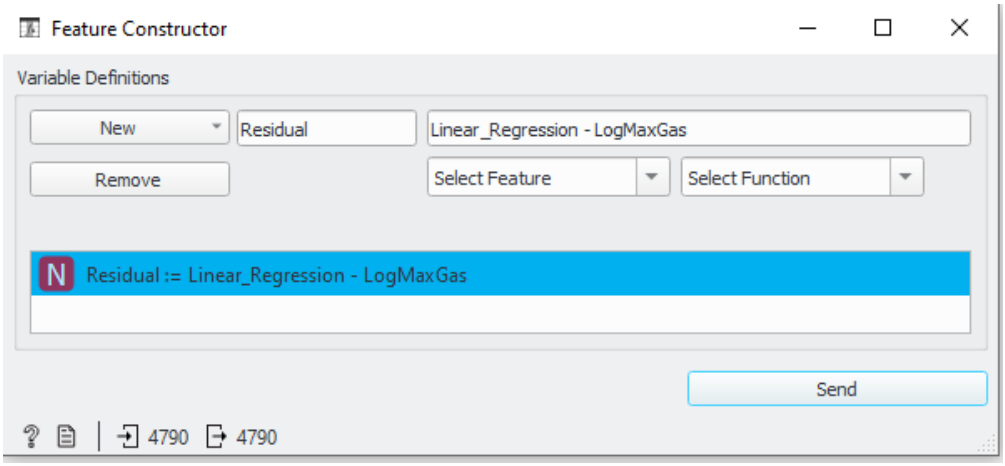

Open and visualize the residual plot to see the effect of the Log transformation of Max_Gas target variable:

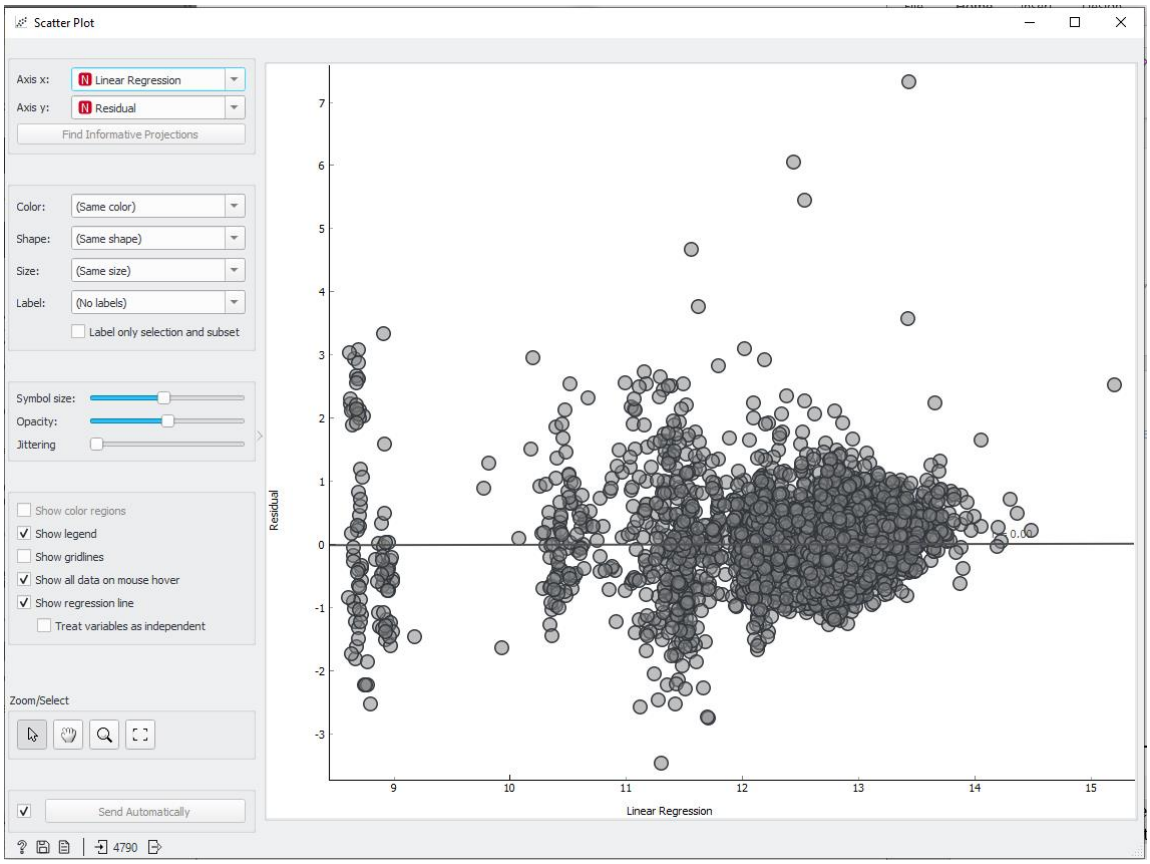

13. Open **Test and Score** widget and complete the table below. Notice the improvement of the model performance.

Model	RMSE	MAE	R2
Linear Regression, LogMaxGas (Cross Validation, 10 folds)	0.590	0.428	0.655

14. If you do a prediction using this model, what is an important consideration before using the predicted results?

RMS value not the same as before because now a log-transformed variable, thus cannot compare the errors the same way (numerically not the same); but the R2 (r-square) should give an indication that the model improved significantly